

Bayesian statistics with R

8. Heterogeneity and multilevel models (aka mixed models)

Olivier Gimenez

April 2022

Multilevel (aka mixed-effect) models

What are multilevel models?

- Multilevel models include both fixed and random effects.

What are multilevel models?

- Multilevel models include both fixed and random effects.
- Random effects are statistical parameters that attempt to **explain noise caused by clusters** of the population you are trying to model.

What are multilevel models?

- Multilevel models include both fixed and random effects.
- Random effects are statistical parameters that attempt to **explain noise caused by clusters** of the population you are trying to model.
- A multilevel model assumes that the dataset being analysed consists of **a hierarchy of different populations** whose differences relate to that hierarchy.

What are multilevel models?

- Multilevel models include both fixed and random effects.
- Random effects are statistical parameters that attempt to **explain noise caused by clusters** of the population you are trying to model.
- A multilevel model assumes that the dataset being analysed consists of **a hierarchy of different populations** whose differences relate to that hierarchy.
- Measurement that come **in clusters** or groups.

What are multilevel models?

- Multilevel models include both fixed and random effects.
- Random effects are statistical parameters that attempt to **explain noise caused by clusters** of the population you are trying to model.
- A multilevel model assumes that the dataset being analysed consists of **a hierarchy of different populations** whose differences relate to that hierarchy.
- Measurement that come **in clusters** or groups.
- Come up with examples of clusters or groups.

Clusters might be:

- Classrooms within schools
- Students within classrooms
- Chapters within books
- Individuals within populations
- Populations within species
- Trajectories within individuals
- Fishes within tanks
- Frogs within ponds
- PhD applicants in doctoral schools
- Nations in continents
- Sex or age are not clusters per se (if we were to sample again, we would take the same levels, e.g. male/female and young/old)

Why do we need multilevel models?

- Model the clustering itself.

Why do we need multilevel models?

- Model the clustering itself.
- Interested in variance components (environmental vs. genetic variance).

Why do we need multilevel models?

- Model the clustering itself.
- Interested in variance components (environmental vs. genetic variance).
- Control for bias due to pseudoreplication (time, space, individual).

McElreath's explanation of multilevel models

- Fixed-effect models have amnesia.

McElreath's explanation of multilevel models

- Fixed-effect models have amnesia.
- Every new cluster (individual, species, classroom) is a new world.

McElreath's explanation of multilevel models

- Fixed-effect models have amnesia.
- Every new cluster (individual, species, classroom) is a new world.
- No information passed among clusters.

McElreath's explanation of multilevel models

- Fixed-effect models have amnesia.
- Every new cluster (individual, species, classroom) is a new world.
- No information passed among clusters.
- Multilevel models remember and pool information. They have memory.

McElreath's explanation of multilevel models

- Fixed-effect models have amnesia.
- Every new cluster (individual, species, classroom) is a new world.
- No information passed among clusters.
- Multilevel models remember and pool information. They have memory.
- Properties of clusters come from a population.

McElreath's explanation of multilevel models


- Fixed-effect models have amnesia.
- Every new cluster (individual, species, classroom) is a new world.
- No information passed among clusters.
- Multilevel models remember and pool information. They have memory.
- Properties of clusters come from a population.
- If previous clusters improve your guess about a new cluster, you want to use pooling.

Plant experiment in the field at CEFE



Courtesy of Pr Eleni Kazakou

Number of grains per species (cluster) as a function of biomass



GLM with complete pooling

$$Y_i \sim \text{Distribution}(\text{mean}_i) \quad [\text{likelihood}]$$

$$\text{link}(\text{mean})_i = \alpha + \beta x_i \quad [\text{linear model}]$$

$$\alpha \sim \text{to be determined} \quad [\text{prior for intercept}]$$

$$\beta \sim \text{to be determined} \quad [\text{prior for slope}]$$

Model with complete pooling. All clusters the same.

GLM with no pooling

$$Y_i \sim \text{Distribution}(\text{mean}_i) \quad [\text{likelihood}]$$

$$\text{link}(\text{mean})_i = \alpha_{\text{CLUSTER}[i]} + \beta x_i \quad [\text{linear model}]$$

$$\alpha_j \sim \text{to be determined} \quad [\text{prior for intercept}]$$

$$\beta \sim \text{to be determined} \quad [\text{prior for slope}]$$

Model with no pooling. All clusters unrelated (fixed effect).

GLMM or GLM with partial pooling

$Y_i \sim \text{Distribution}(\text{mean}_i)$	[likelihood]
$\text{link}(\text{mean})_i = \alpha_{\text{CLUSTER}[i]} + \beta x_i$	[linear model]
$\alpha_j \sim \text{Normal}(\bar{\alpha}, \sigma)$	[prior for varying intercepts]
$\bar{\alpha} \sim \text{to be determined}$	[prior for population mean]
$\sigma \sim \text{to be determined}$	[prior for standard deviation]
$\beta \sim \text{to be determined}$	[prior for slope]

Model with partial pooling. Clusters are somehow related (random effect).

Back to the plant example

Model with complete pooling (all species are the same)

$n\text{seeds}_i \sim \text{Normal}(\mu_i, \sigma^2)$ [likelihood]

$\mu_i = \alpha + \beta \text{ biomass}_i$ [linear model]

$\alpha \sim \text{Normal}(0, 1000)$ [prior for intercept]

$\beta \sim \text{Normal}(0, 1000)$ [prior for slope]

$\sigma \sim \text{Uniform}(0, 100)$ [prior for standard deviation]

Read in and manipulate data

```
# read in data
VMG <- read_csv2(here::here("slides", "dat", "VMG.csv")) %>%
  mutate(Sp = as_factor(Sp), Vm = as.numeric(Vm))

# nb of seeds
y <- log(VMG$NGrTotest)

# biomass
x <- VMG$Vm
x <- (x - mean(x))/sd(x)

# species name
Sp <- VMG$Sp

# species label
species <- as.numeric(Sp)

# species name
nbspecies <- length(levels(Sp))
```

Specify the model in Jags

```
model <-  
paste("  
model{  
  for(i in 1:n){  
    y[i] ~ dnorm(mu[i], tau.y)  
    mu[i] <- a + b * x[i]  
  }  
  tau.y <- 1 / (sigma.y * sigma.y)  
  sigma.y ~ dunif(0,100)  
  a ~ dnorm(0,0.001)  
  b ~ dnorm(0,0.001)  
}  
")  
writeLines(model,here::here("slides","code","completepooling.bug"))
```

Prepare ingredients for running Jags

data

```
allom.data <- list(y = y, n = n, x = x)
```

initial values

```
init1 <- list(a=rnorm(1), b=rnorm(1),sigma.y=runif(1))
```

```
init2 <- list(a=rnorm(1), b=rnorm(1),sigma.y=runif(1))
```

```
inits <- list(init1,init2)
```

parameters to be estimated

```
allom.parameters <- c("a", "b", "sigma.y")
```

Run Jags

```
allom.1 <- jags(allom.data,  
               inits,  
               allom.parameters,  
               n.iter = 2500,  
               model.file = here::here("slides", "code", "completepooling.br",  
               n.chains = 2,  
               n.burn = 1000)  
  
#> Compiling model graph  
#>   Resolving undeclared variables  
#>   Allocating nodes  
#> Graph information:  
#>   Observed stochastic nodes: 488  
#>   Unobserved stochastic nodes: 3  
#>   Total graph size: 1956
```

Display results

```
allom.1
```

```
#> Inference for Bugs model at "/Users/oliviergimenez/Dropbox/OG/GITHUB/bugs" using JAGS.
```

```
#> 2 chains, each with 2500 iterations (first 1000 discarded)
```

```
#> n.sims = 3000 iterations saved
```

```
#>          mu.vect sd.vect      2.5%      25%      50%      75%      97.5%
```

```
#> a          13.927   0.464   13.009   13.615   13.929   14.245   14.825
```

```
#> b           3.581   0.473    2.662    3.269    3.582    3.898    4.491
```

```
#> sigma.y     10.431   0.346    9.803   10.205   10.417   10.651   11.156
```

```
#> deviance 3672.061    2.533 3669.221 3670.245 3671.432 3673.199 3678.795
```

```
#>          n.eff
```

```
#> a          3000
```

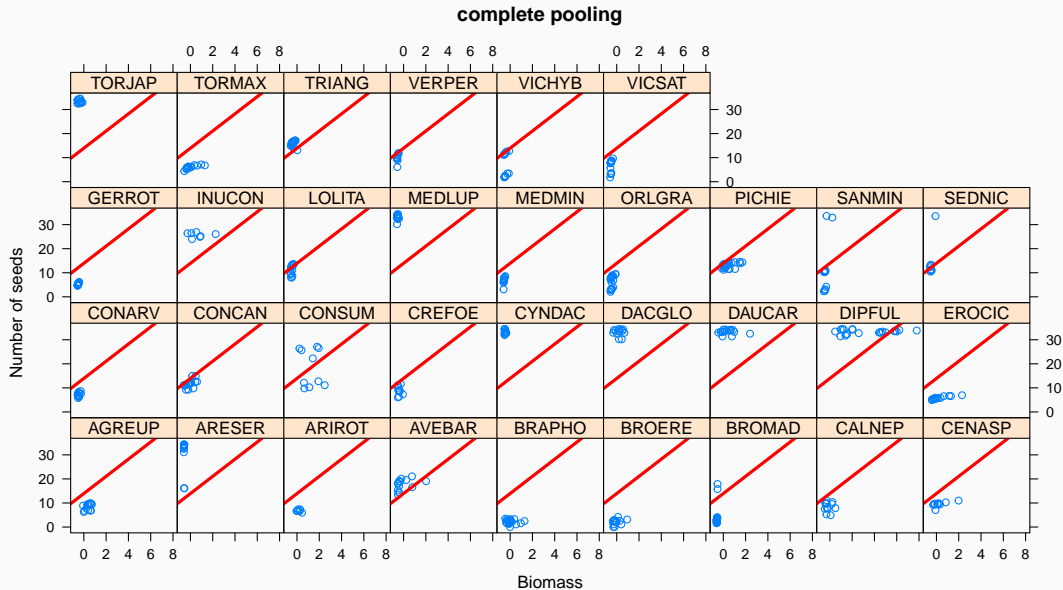
```
#> b          3000
```

```
#> sigma.y     3000
```

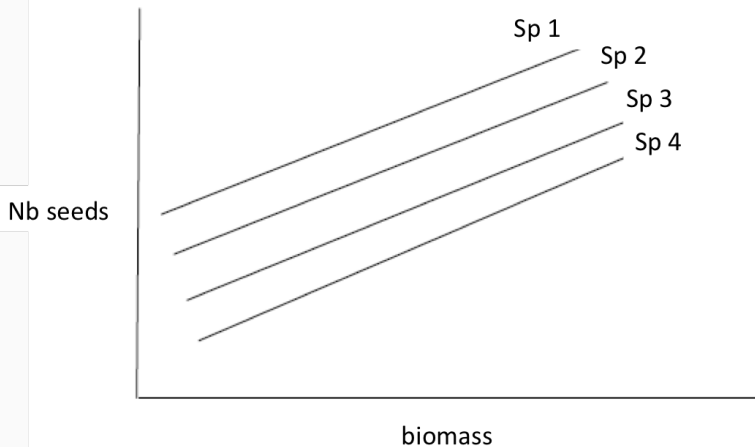
```
#> deviance     3000
```

Compare with Frequentist approach

```
freq_lm <- lm(y ~ x, data = allom.data)
freq_lm
#>
#> Call:
#> lm(formula = y ~ x, data = allom.data)
#>
#> Coefficients:
#> (Intercept)          x
#>      13.927       3.578
```



Model with partial pooling (species random effect)



Model with partial pooling (all species related in some way)

$n\text{seeds}_i \sim \text{Normal}(\mu_i, \sigma^2)$ [likelihood]

$\mu_i = \alpha_{\text{species}[i]} + \beta \text{biomass}_i$ [linear model]

$\alpha_j \sim \text{Normal}(\bar{\alpha}, \sigma_\alpha)$ [prior for varying intercepts]

$\bar{\alpha} \sim \text{Normal}(0, 1000)$ [prior for population mean]

$\sigma_\alpha \sim \text{Uniform}(0, 100)$ [prior for σ_α]

$\beta \sim \text{Normal}(0, 1000)$ [prior for slope]

$\sigma \sim \text{Uniform}(0, 100)$ [prior for σ]

Implementation in Jags

```
model <- paste("
model {
  for (i in 1:n){
    y[i] ~ dnorm(mu[i], tau.y)
    mu[i] <- a[species[i]] + b * x[i]
  }
  tau.y <- 1/ (sigma.y * sigma.y)
  sigma.y ~ dunif(0, 100)
  for (j in 1:nbspecies){
    a[j] ~ dnorm(mu.a, tau.a)
  }
  mu.a ~ dnorm(0, 0.001)
  tau.a <- 1/(sigma.a * sigma.a)
  sigma.a ~ dunif(0, 100)
```

Prepare ingredients for running Jags

```
allom.data <- list(n = n,  
                  nbspecies = nbspecies,  
                  x = x,  
                  y = y,  
                  species = species)  
  
init1 <- list(a = rnorm(nbspecies), b = rnorm(1), mu.a = rnorm(1),  
             sigma.y = runif(1), sigma.a=runif(1))  
  
init2 <- list(a = rnorm(nbspecies), b = rnorm(1), mu.a = rnorm(1),  
             sigma.y = runif(1), sigma.a = runif(1))  
  
inits <- list(init1,init2)  
  
allom.parameters <- c("b", "mu.a","sigma.y", "sigma.a")
```

Run Jags

```
allom.2 <- jags(allom.data,  
               inits,  
               allom.parameters,  
               n.iter = 2500,  
               model.file = here::here("slides", "code", "varint.bug"),  
               n.chains = 2,  
               n.burn = 1000)  
  
#> Compiling model graph  
#>   Resolving undeclared variables  
#>   Allocating nodes  
#> Graph information:  
#>   Observed stochastic nodes: 488  
#>   Unobserved stochastic nodes: 37  
#>   Total graph size: 2481
```

Display results

allom.2

#> Inference for Bugs model at "/Users/oliviergimenez/Dropbox/OG/GITHUB/bugs2/bugs2.Rout" using

#> 2 chains, each with 2500 iterations (first 1000 discarded)

#> n.sims = 3000 iterations saved

#>	mu.vect	sd.vect	2.5%	25%	50%	75%	97.5%
----	---------	---------	------	-----	-----	-----	-------

#> b	0.480	0.238	0.015	0.318	0.479	0.638	0.938
------	-------	-------	-------	-------	-------	-------	-------

#> mu.a	14.502	1.993	10.708	13.178	14.518	15.828	18.436
---------	--------	-------	--------	--------	--------	--------	--------

#> sigma.a	11.096	1.458	8.605	10.089	10.962	11.957	14.331
------------	--------	-------	-------	--------	--------	--------	--------

#> sigma.y	3.065	0.099	2.878	2.997	3.063	3.132	3.258
------------	-------	-------	-------	-------	-------	-------	-------

#> deviance	2478.046	8.542	2463.245	2471.869	2477.460	2483.513	2496.398
-------------	----------	-------	----------	----------	----------	----------	----------

#>	n.eff
----	-------

#> b	3000
------	------


#> mu.a	3000
---------	------

#> sigma.a	1200
------------	------

Compare with Frequentist approach

```
library(lme4)
freq_lmm <- lmer(y ~ x + (1 | species), allom.data, REML = FALSE)
freq_lmm
#> Linear mixed model fit by maximum likelihood ['lmerMod']
#> Formula: y ~ x + (1 | species)
#> Data: allom.data
#> AIC      BIC    logLik deviance df.resid
#> 2652.606 2669.368 -1322.303 2644.606      484
#> Random effects:
#> Groups   Name      Std.Dev.
#> species (Intercept) 10.472
#> Residual              3.058
#> Number of obs: 488, groups: species, 33
#> Fixed Effects:
```

Compare complete pooling vs partial pooling



Model with no pooling (all species unrelated)

$n\text{seeds}_i \sim \text{Normal}(\mu_i, \sigma^2)$ [likelihood]

$\mu_i = \alpha_{\text{species}[i]} + \beta \text{biomass}_i$ [linear model]

$\alpha_j \sim \text{Normal}(0, 1000)$ [prior for intercepts]

$\beta \sim \text{Normal}(0, 1000)$ [prior for slope]

$\sigma \sim \text{Uniform}(0, 100)$ [prior for σ]

Implementation in Jags

```
model <- paste("
model {
  for (i in 1:n){
    y[i] ~ dnorm (mu[i], tau.y)
    mu[i] <- a[species[i]] + b * x[i]
  }
  tau.y <- 1 / (sigma.y * sigma.y)
  sigma.y ~ dunif(0, 100)
  for (j in 1:nbspecies){
    a[j] ~ dnorm(0, 0.001)
  }
  b ~ dnorm(0,0.1)
}")

writeLines(model here::here("slides" "code" "nonpooling bug"))
```

Prepare ingredients

```
allom.data <- list(n = n, nbSpecies = nbSpecies, x = x, y = y, species = sp)
init1 <- list(a = rnorm(nbSpecies), b = rnorm(1), sigma.y = runif(1))
init2 <- list(a = rnorm(nbSpecies), b = rnorm(1), sigma.y = runif(1))
inits<-list(init1, init2)
allom.parameters <- c("a","b","sigma.y")
```

Run JAGS

```
allom.3 <- jags(data = allom.data,  
               inits = inits,  
               parameters.to.save = allom.parameters,  
               n.iter = 2500,  
               model.file = here::here("slides", "code", "nopooling.bug"),  
               n.chains = 2,  
               n.burn = 1000)  
  
#> Compiling model graph  
#>   Resolving undeclared variables  
#>   Allocating nodes  
#> Graph information:  
#>   Observed stochastic nodes: 488  
#>   Unobserved stochastic nodes: 35  
#>   Total graph size: 2481
```

Display results


```
allom.3$BUGSoutput$summary[c(1:4, 32:33, 34), -c(4,6)]
```

#>		<i>mean</i>	<i>sd</i>	<i>2.5%</i>	<i>50%</i>	<i>97.5%</i>	<i>Rhat</i>	<i>n.e.</i>
#> a[1]	8.1659636	0.8157488	6.590577	8.1798267	9.8082250	1.000858	30	
#> a[2]	30.7338212	0.8846737	29.017299	30.7379259	32.4311725	1.001322	23	
#> a[3]	6.6512875	1.1509644	4.414152	6.6473330	8.9381659	1.004078	4	
#> a[4]	17.6301926	0.7809627	16.074789	17.6426587	19.1633082	1.000666	30	
#> a[32]	6.3626714	0.7948928	4.799545	6.3630696	7.9474871	1.003034	6	
#> a[33]	6.5942650	0.8035074	5.034882	6.5921967	8.1412549	1.000843	30	
#> b	0.4367871	0.2469271	-0.066436	0.4359984	0.9067317	1.000950	30	

Compare with Frequentist approach

```
lm(y ~ -1 + as.factor(species) + x, data = allom.data) %>%  
  broom::tidy() %>%  
  slice(c(1:4, 32:33, 34))  
#> # A tibble: 7 x 5  
#>   term                estimate std.error statistic    p.value  
#>   <chr>              <dbl>      <dbl>      <dbl>    <dbl>  
#> 1 as.factor(species)1      8.17      0.824      9.92 3.92e- 21  
#> 2 as.factor(species)2     30.8      0.895     34.4 1.67e-128  
#> 3 as.factor(species)3      6.67      1.16      5.76 1.56e- 8  
#> 4 as.factor(species)4     17.6      0.791     22.3 5.32e- 75  
#> 5 as.factor(species)32      6.38      0.797      8.01 9.95e- 15  
#> 6 as.factor(species)33      6.63      0.800      8.29 1.33e- 15  
#> 7 x                0.441      0.243      1.81 7.06e- 2
```

Compare complete pooling vs partial pooling vs no pooling



Bonus: Model with varying intercept and varying slope

Code: part 1

```
model <-  
paste("  
# varying-intercept, varying-slope allometry model  
# with Vm as a species predictor  
  
model {  
  for (i in 1:n){  
    y[i] ~ dnorm (mu[i], tau.y)  
    mu[i] <- a[species[i]] + b[species[i]] * x[i]  
  }  
  
  tau.y <- pow(sigma.y, -2)  
  sigma.y ~ dunif (0, 100)
```


Code: part 2

```
for (j in 1:nbspecies){  
  a[j] ~ dnorm (mu.a, tau.a)  
  b[j] ~ dnorm (mu.b, tau.b)  
}
```

```
mu.a ~ dnorm (0, .001)  
tau.a <- pow(sigma.a, -2)  
sigma.a ~ dunif (0, 100)
```

```
mu.b ~ dnorm (0, .001)  
tau.b <- pow(sigma.b, -2)  
sigma.b ~ dunif (0, 100)
```

```
}
```

Prepare ingredients

```
init1 <- list(a = rnorm(nbspecies), b = rnorm(nbspecies),  
             mu.a = rnorm(1), mu.b = rnorm(1),  
             sigma.y = runif(1), sigma.a = runif(1), sigma.b = runif(1))  
init2 <- list(a = rnorm(nbspecies), b = rnorm(nbspecies),  
             mu.a = rnorm(1), mu.b = rnorm(1),  
             sigma.y = runif(1), sigma.a = runif(1), sigma.b = runif(1))  
inits <- list(init1, init2)  
allom.parameters <- c("a", "b", "mu.a", "mu.b", "sigma.y", "sigma.a", "sigma.b")
```

Run Jags

```
allom.4 <- jags(data = allom.data,  
               inits = inits,  
               parameters.to.save = allom.parameters,  
               n.iter = 2500,  
               model.file = here::here("slides", "code", "varintvarslope.bug"),  
               n.chains = 2,  
               n.burn = 1000)  
  
#> Compiling model graph  
#>   Resolving undeclared variables  
#>   Allocating nodes  
#> Graph information:  
#>   Observed stochastic nodes: 488  
#>   Unobserved stochastic nodes: 71  
#>   Total graph size: 2521
```

Display results

```
round(allom.4$BUGSoutput$summary[c(1:2, 32:33, 34:35, 65:66, 68:72), -c(4,6)], 2)
```

```
#>           mean      sd   2.5%   50% 97.5% Rhat n.eff
#> a[1]         7.63   1.32    4.93    7.63 10.16 1.04     57
#> a[2]        24.20   6.55    7.23   24.94 35.50 1.11     37
#> a[32]         8.57   1.87    4.93    8.59 12.25 1.02    120
#> a[33]        12.97   3.84    4.67   13.12 20.37 1.10     27
#> b[1]          1.91   2.87   -3.58    1.85  7.56 1.03     54
#> b[2]       -10.91  11.36  -40.87   -9.57   8.86 1.12     36
#> b[32]          5.70   4.30   -2.77    5.75 13.87 1.02    130
#> b[33]        12.77   7.39   -2.96   12.93 26.96 1.11     19
#> mu.a         16.70   2.02   12.80   16.67 20.71 1.02     76
#> mu.b          5.13   2.47    0.46    5.00 10.22 1.09     28
#> sigma.a       10.91   1.48    8.46   10.76 14.26 1.01    140
#> sigma.b       12.21   2.68    7.86   11.90 18.43 1.11     7
```

Compare with Frequentist approach

```
freq_lmm2 <- lmer (y ~ x + (1 + x | species), allom.data, REML = FALSE)
freq_lmm2
#> Linear mixed model fit by maximum likelihood ['lmerMod']
#> Formula: y ~ x + (1 + x | species)
#> Data: allom.data
#>      AIC      BIC    logLik deviance df.resid
#> 2609.941 2635.083 -1298.971  2597.941      482
#> Random effects:
#> Groups   Name      Std.Dev. Corr
#> species (Intercept) 10.409
#>          x          11.325   0.22
#> Residual          2.652
#> Number of obs: 488, groups:  species, 33
#> Fixed Effects:
```

Compare with Frequentist approach - with no correlation

```
freq_lmm_wocorr <- lmer(y ~ x + (1 | species) +  
                        (0 + x | species), allom.data, REML = F)  
freq_lmm_wocorr  
#> Linear mixed model fit by maximum likelihood ['lmerMod']  
#> Formula: y ~ x + (1 | species) + (0 + x | species)  
#> Data: allom.data  
#> AIC      BIC    logLik deviance df.resid  
#> 2609.086 2630.037 -1299.543 2599.086      483  
#> Random effects:  
#> Groups   Name      Std.Dev.  
#> species (Intercept) 10.203  
#> species.1 x          10.632  
#> Residual          2.661  
#> Number of obs: 488 groups: species 33
```

Shrinkage results from pooling of information

- Varying effect estimates shrink towards mean ($\bar{\alpha}$).

Shrinkage results from pooling of information

- Varying effect estimates shrink towards mean ($\bar{\alpha}$).
- Avoids underfitting as in complete pooling model (null variance) or overfitting as in no pooling model (infinite variance).

Shrinkage results from pooling of information

- Varying effect estimates shrink towards mean ($\bar{\alpha}$).
- Avoids underfitting as in complete pooling model (null variance) or overfitting as in no pooling model (infinite variance).
- Varying effects: adaptive regularization through cluster variance estimation.

Shrinkage results from pooling of information

- Varying effect estimates shrink towards mean ($\bar{\alpha}$).
- Avoids underfitting as in complete pooling model (null variance) or overfitting as in no pooling model (infinite variance).
- Varying effects: adaptive regularization through cluster variance estimation.
- Further from mean, more shrinkage.

Shrinkage results from pooling of information

- Varying effect estimates shrink towards mean ($\bar{\alpha}$).
- Avoids underfitting as in complete pooling model (null variance) or overfitting as in no pooling model (infinite variance).
- Varying effects: adaptive regularization through cluster variance estimation.
- Further from mean, more shrinkage.
- Fewer data in cluster, more shrinkage.

Multilevel models are awesome!

Multilevel models in a nutshell

- **Shrinkage via pooling is desirable.** The no-pooling model overstates variation among clusters and makes the individual clusters look more different than they are (overfitting). The complete-pooling model simply ignores the variation among clusters (underfitting).

Multilevel models in a nutshell

- **Shrinkage via pooling is desirable.** The no-pooling model overstates variation among clusters and makes the individual clusters look more different than they are (overfitting). The complete-pooling model simply ignores the variation among clusters (underfitting).
- We can **generalize to a wider population**. Is there an allometry relationship between number of seeds and biomass?

Multilevel models in a nutshell

- **Shrinkage via pooling is desirable.** The no-pooling model overstates variation among clusters and makes the individual clusters look more different than they are (overfitting). The complete-pooling model simply ignores the variation among clusters (underfitting).
- We can **generalize to a wider population**. Is there an allometry relationship between number of seeds and biomass?
- We may consider **varying slopes**. We'd need to deal with correlations between intercept and slope random effects. Open a whole new world with spatial (or time) autocorrelation, phylogenetic regressions, quantitative genetics, network models.

Multilevel models in a nutshell

- **Shrinkage via pooling is desirable.** The no-pooling model overstates variation among clusters and makes the individual clusters look more different than they are (overfitting). The complete-pooling model simply ignores the variation among clusters (underfitting).
- We can **generalize to a wider population**. Is there an allometry relationship between number of seeds and biomass?
- We may consider **varying slopes**. We'd need to deal with correlations between intercept and slope random effects. Open a whole new world with spatial (or time) autocorrelation, phylogenetic regressions, quantitative genetics, network models.
- We may **include predictors at the cluster level**. Imagine we know something about functional traits, and wish to determine whether some species-to-species variation in the allometry relationship is explained by these traits.

Your turn: Practical 8

Conclusions

Take-home messages about Bayesian statistics

- Frees the modeler in you (M. Kéry)
 - Uses probability to quantify uncertainty for everything (propagation of uncertainty).
 - Allows use of prior information ('better' estimates).
 - Can fit complex (hierarchical) models with same MCMC algorithms.


Take-home messages about Bayesian statistics

- Frees the modeler in you (M. Kéry)
 - Uses probability to quantify uncertainty for everything (propagation of uncertainty).
 - Allows use of prior information ('better' estimates).
 - Can fit complex (hierarchical) models with same MCMC algorithms.
- With great tools come great responsibilities
 - Checking convergence is painful.
 - Specifying priors might be tricky.
 - Model adequacy should be checked (posterior predictive checks - not covered).
 - Computational burden can be high (see function `R2jags::jags.parallel()` and package 'jagsUI'.

Take-home messages about Bayesian statistics

- Frees the modeler in you (M. Kéry)
 - Uses probability to quantify uncertainty for everything (propagation of uncertainty).
 - Allows use of prior information ('better' estimates).
 - Can fit complex (hierarchical) models with same MCMC algorithms.
- With great tools come great responsibilities
 - Checking convergence is painful.
 - Specifying priors might be tricky.
 - Model adequacy should be checked (posterior predictive checks - not covered).
 - Computational burden can be high (see function `R2jags::jags.parallel()` and package '`jagsUI`').
- So what?
 - Make an informed and pragmatic choice.
 - Are you after complexity, speed, uncertainties, etc?
 - Talk to colleagues.

(YET ANOTHER) HISTORY OF LIFE AS WE KNOW IT...



HOMO
APRIORIUS

HOMO
PRAGMATICUS

HOMO
FREQUENTISTUS

HOMO
SAPIENS

HOMO
BAYESIANIS

Why become a bayesian? Ask twitter!



Chelsea Parlett-Pelleriti
@ChelseaParlett

Why did you become a Bayesian, wrong answers only

[Traduire le Tweet](#)



Your turn: Practical 9
