

Bayesian statistics with R

7. Contrast scientific hypotheses with model selection

Olivier Gimenez

November 2021

Model selection

How to select a best model?

- Is there any effect of rain or temperature or both on breeding success?

How to select a best model?

- Is there any effect of rain or temperature or both on breeding success?
- The proportion of explained variance R^2 is problematic, because the more variables you have, the bigger R^2 is.

How to select a best model?

- Is there any effect of rain or temperature or both on breeding success?
- The proportion of explained variance R^2 is problematic, because the more variables you have, the bigger R^2 is.
- Idea: **penalize models with too many parameters.**

Akaike information criterion (AIC)

$$AIC = -2 \log(L(\hat{\theta}_1, \dots, \hat{\theta}_K)) + 2K$$

with L the likelihood and K the number of parameters θ_i .

Akaike information criterion (AIC)

$$\text{AIC} = -2 \log(L(\hat{\theta}_1, \dots, \hat{\theta}_K)) + 2K$$

A measure of goodness-of-fit of the model to the data: the more parameters you have, the smaller the deviance is (or the bigger the likelihood is).

Akaike information criterion (AIC)

$$\text{AIC} = -2 \log(L(\hat{\theta}_1, \dots, \hat{\theta}_K)) + 2K$$

A **penalty**: twice the number of parameters K

Akaike information criterion (AIC)

- AIC makes the balance between *quality of fit* and *complexity* of a model.

Akaike information criterion (AIC)

- AIC makes the balance between *quality of fit* and *complexity* of a model.
- Best model is the one with lowest AIC value.

Akaike information criterion (AIC)

- AIC makes the balance between *quality of fit* and *complexity* of a model.
- Best model is the one with lowest AIC value.
- Two models are difficult to distinguish if $\Delta AIC < 2$.

- Watanabe-Akaike (Widely-Applicable) Information Criteria or WAIC:

$$\text{WAIC} = -2 \sum_{i=1}^n \log E[\Pr(y_i \mid \theta)] + 2p_{\text{WAIC}}$$

- where $E[p(y_i \mid \theta)]$ is the posterior mean of the likelihood evaluated pointwise at each i th observation.
- p_{WAIC} is a penalty computed using the posterior variance of the likelihood.
- More in this video <https://www.youtube.com/watch?v=vSjL2Zc-gEQ> by McElreath.
- Relatively new and not yet available in Jags in routine.

WAIC in Jags

```
# calculate wAIC with JAGS  
# https://sourceforge.net/p/mcmc-jags/discussion/610036/thread/8211df61/#ea5c  
samples <- jags.samples(storks$model, c("WAIC", "deviance"), type = "mean",  
                        n.iter = 2000,  
                        n.burnin = 1000,  
                        n.thin = 1)
```

WAIC in Jags

```
samples$p_waic <- samples$WAIC
samples$waic <- samples$deviance + samples$p_waic
tmp <- sapply(samples, sum)
waic <- round(c(waic = tmp[["waic"]], p_waic = tmp[["p_waic"]]),1)
waic
#>    waic p_waic
#> 217.4   12.8
```

Further reading

- Hooten, M.B. and Hobbs, N.T. (2015), A guide to Bayesian model selection for ecologists. Ecological Monographs, 85: 3-28. <https://doi.org/10.1890/14-0661.1>
- Conn, P.B., Johnson, D.S., Williams, P.J., Melin, S.R. and Hooten, M.B. (2018), A guide to Bayesian model checking for ecologists. Ecol Monogr, 88: 526-542. <https://doi.org/10.1002/ecm.1314>

Your turn: Practical 7
